

# Handling missing data

```
In [ ]: import numpy as np
import pandas as pd
```

## Pandas Utility functions:

Similarly to numpy, pandas also has a few utility functions to identify and detect null values:

```
In [ ]: print(pd.isnull(np.nan), pd.isnull(None), pd.isna(np.nan), pd.isna(None))
```

True True True True

the opposite also exists:

```
In [ ]: print(pd.notnull(None), pd.notnull(np.nan), pd.notna(None), pd.notna(np.nan), pd.notna(1))
```

False False False False True

These also work with series and Dataframes

```
In [ ]: print(pd.isnull(pd.Series([1, np.nan, 7])), '\n', pd.notnull(pd.Series([1, np.nan, 7])))
```

```
0    False
1     True
2    False
dtype: bool
0     True
1    False
2     True
dtype: bool
```

```
In [ ]: pd.isnull(pd.DataFrame({
    'Column A': [1, np.nan, 7],
    'Column B': [np.nan, 2, 3],
    'Column C': [np.nan, 2, np.nan]
})))
```

```
Out[ ]:   Column A  Column B  Column C
0      False      True      True
1       True     False     False
2      False     False      True
```

## Filtering missing data

```
In [ ]: s = pd.Series([1, 2, 3, np.nan, np.nan, 4])
s
```

```
Out[ ]: 0    1.0
        1    2.0
        2    3.0
        3    NaN
        4    NaN
        5    4.0
        dtype: float64
```

```
In [ ]: print("Table with boolean for not null: ",'\n',pd.notnull(s),'\n','Table with boolean for null: ','\n',pd.isnull(s),'\n')

Table with boolean for not null:
0    True
1    True
2    True
3    False
4    False
5    True
dtype: bool
Table with boolean for null:
0    False
1    False
2    False
3    True
4    True
5    False
dtype: bool
```

```
In [ ]: print("number of null entries: ",pd.isnull(s).sum(),'\n','number of allowed entries: ',pd.notnull(s).sum(),'\n')

number of null entries: 2
number of allowed entries: 4
```

```
In [ ]: s[pd.notnull(s)]
```

```
Out[ ]: 0    1.0
        1    2.0
        2    3.0
        5    4.0
        dtype: float64
```

Both isnull and notnull are also functions of dataframes and series, hence we can use them directly!

```
In [ ]: s[s.notnull()]
```

```
Out[ ]: 0    1.0
        1    2.0
        2    3.0
        5    4.0
        dtype: float64
```

## Dropping null values from dataframes:

```
In [ ]: df = pd.DataFrame({
        'Column A': [1, np.nan, 30, np.nan],
        'Column B': [2, 8, 31, np.nan],
        'Column C': [np.nan, 9, 32, 100],
        'Column D': [5, 8, 34, 110],
    })
df
```

```
Out[ ]:
```

|   | Column A | Column B | Column C | Column D |
|---|----------|----------|----------|----------|
| 0 | 1.0      | 2.0      | NaN      | 5        |
| 1 | NaN      | 8.0      | 9.0      | 8        |
| 2 | 30.0     | 31.0     | 32.0     | 34       |
| 3 | NaN      | NaN      | 100.0    | 110      |

```
In [ ]: df.isnull()
```

```
Out[ ]:
```

|   | Column A | Column B | Column C | Column D |
|---|----------|----------|----------|----------|
| 0 | False    | False    | True     | False    |
| 1 | True     | False    | False    | False    |
| 2 | False    | False    | False    | False    |
| 3 | True     | True     | False    | False    |

```
In [ ]: df=df.rename(columns={'Column A':'a','Column B':'b','Column C':'c','Column D':'d'})
```

```
In [ ]: df
```

```
Out[ ]:
```

|   | a    | b    | c     | d   |
|---|------|------|-------|-----|
| 0 | 1.0  | 2.0  | NaN   | 5   |
| 1 | NaN  | 8.0  | 9.0   | 8   |
| 2 | 30.0 | 31.0 | 32.0  | 34  |
| 3 | NaN  | NaN  | 100.0 | 110 |

```
In [ ]: df.dropna() #this ends up dropping all rows with any *null* value
```

```
Out[ ]:
```

|   | a    | b    | c    | d  |
|---|------|------|------|----|
| 2 | 30.0 | 31.0 | 32.0 | 34 |

```
In [ ]: df.dropna(axis='columns') #removes all columns with any null value
```

```
Out[ ]:
```

|   | d   |
|---|-----|
| 0 | 5   |
| 1 | 8   |
| 2 | 34  |
| 3 | 110 |

```
In [ ]: df2 = pd.DataFrame({
    'Column A': [1, np.nan, 30],
    'Column B': [2, np.nan, 31],
    'Column C': [np.nan, np.nan, 100]
})
df2
```

```
Out[ ]:
```

|   | Column A | Column B | Column C |
|---|----------|----------|----------|
| 0 | 1.0      | 2.0      | NaN      |
| 1 | NaN      | NaN      | NaN      |
| 2 | 30.0     | 31.0     | 100.0    |

now we'll see "how"-> 'any' means that it'll drop all rows/columns with a null value 'all' means that it'll drop all rows/columns with ALL null values

```
In [ ]: df2.dropna(how='any')
```

```
Out[ ]:
```

|   | Column A | Column B | Column C |
|---|----------|----------|----------|
| 2 | 30.0     | 31.0     | 100.0    |

```
In [ ]: df2.dropna(how='all')
```

```
Out[ ]:
```

|   | Column A | Column B | Column C |
|---|----------|----------|----------|
| 0 | 1.0      | 2.0      | NaN      |
| 2 | 30.0     | 31.0     | 100.0    |

```
In [ ]: df2.dropna(thresh=3) #how many non-na values to keep by default looks at indexes (s
```

```
Out[ ]:
```

|   | Column A | Column B | Column C |
|---|----------|----------|----------|
| 2 | 30.0     | 31.0     | 100.0    |

```
In [ ]: df.dropna(thresh=3, axis='columns')
```

```
Out[ ]:
```

|   | b    | c     | d   |
|---|------|-------|-----|
| 0 | 2.0  | NaN   | 5   |
| 1 | 8.0  | 9.0   | 8   |
| 2 | 31.0 | 32.0  | 34  |
| 3 | NaN  | 100.0 | 110 |

## Filling out null values:

Sometimes instead than dropping the null values, we might need to replace them with some other value. This highly depends on your context and the dataset you're currently working. Sometimes a nan can be replaced with a 0, sometimes it can be replaced with the mean of the sample, and some other times you can take the closest value. Again, it depends on the context. We'll show you the different methods and mechanisms and you can then apply them to your own problem.

```
In [ ]: s
```

```
Out[ ]: 0    1.0
        1    2.0
        2    3.0
        3    NaN
        4    NaN
        5    4.0
        dtype: float64
```

```
In [ ]: s.fillna(0) #replaces all Nan values with 0
```

```
Out[ ]: 0    1.0
        1    2.0
        2    3.0
        3    0.0
        4    0.0
        5    4.0
        dtype: float64
```

```
In [ ]: s.fillna(s.mean())
```

```
Out[ ]: 0    1.0
        1    2.0
        2    3.0
        3    2.5
        4    2.5
        5    4.0
        dtype: float64
```

The "method argument is used to fill null values with other values close to that null one:

```
In [ ]: s.fillna(method='ffill') #f forward
```

```
Out[ ]: 0    1.0
        1    2.0
        2    3.0
        3    3.0
        4    3.0
        5    4.0
        dtype: float64
```

```
In [ ]: s.fillna(method='bfill') # b backwards
```

```
Out[ ]: 0    1.0
        1    2.0
        2    3.0
        3    4.0
        4    4.0
        5    4.0
        dtype: float64
```

the following method always leaves null values at an extreme end (maybe even both!)

```
In [ ]: pd.Series([np.nan, 3, np.nan, 9]).fillna(method='ffill')
```

```
Out[ ]: 0    NaN
        1    3.0
        2    3.0
        3    9.0
        dtype: float64
```

## Filna also works with DataFrames:

The fillna method also works on DataFrames, and it works similarly. The main differences are that you can specify the axis (as usual, rows or columns) to use to fill the values (specially for methods) and that you have more control on the values passed:

```
In [ ]: df
```

```
Out[ ]:
```

|   | a    | b    | c     | d   |
|---|------|------|-------|-----|
| 0 | 1.0  | 2.0  | NaN   | 5   |
| 1 | NaN  | 8.0  | 9.0   | 8   |
| 2 | 30.0 | 31.0 | 32.0  | 34  |
| 3 | NaN  | NaN  | 100.0 | 110 |

```
In [ ]: df.fillna({'a': 0, 'b': 99, 'c': df['c'].mean()})
```

```
Out[ ]:
```

|   | a    | b    | c     | d   |
|---|------|------|-------|-----|
| 0 | 1.0  | 2.0  | 47.0  | 5   |
| 1 | 0.0  | 8.0  | 9.0   | 8   |
| 2 | 30.0 | 31.0 | 32.0  | 34  |
| 3 | 0.0  | 99.0 | 100.0 | 110 |

```
In [ ]: df.fillna(method='ffill', axis=0) #you can see here how it forgot an extreme!  
#AXIS=0 IS COLUMNS AND AXIS=1 IS ROW BASED.
```

```
Out[ ]:
```

|   | a    | b    | c     | d   |
|---|------|------|-------|-----|
| 0 | 1.0  | 2.0  | NaN   | 5   |
| 1 | 1.0  | 8.0  | 9.0   | 8   |
| 2 | 30.0 | 31.0 | 32.0  | 34  |
| 3 | 30.0 | 31.0 | 100.0 | 110 |

## Checking if there are any NAs:

The question is: Does this Series or DataFrame contain any missing value? The answer should be yes or no: True or False. How can you verify it?

```
In [ ]: s.dropna().count()
```

```
Out[ ]: 4
```

```
In [ ]: missing_values = len(s.dropna()) != len(s)  
missing_values
```

```
Out[ ]: True
```

## Pythonic solutions:

The methods `any` and `all` check if either there's any true value in a Series or all the values are True.

```
In [ ]: print(pd.Series([True, False, False]).any(), pd.Series([True, True, True]).all())
```

```
True True
```

```
In [ ]: pd.Series([1, 2]).isnull().any()
```

```
Out[ ]: False
```

```
In [ ]: s.isnull().values
```

```
Out[ ]: array([False, False, False,  True,  True, False])
```

```
In [ ]: s.isnull().values.any()
```

```
Out[ ]: True
```